

# Generative adversarial networks



Ian  
Goodfellow



Jean  
Pouget-Abadie



Mehdi  
Mirza



Bing  
Xu



David  
Warde-Farley



Sherjil  
Ozair




Aaron  
Courville



Yoshua  
Bengio


# Discriminative deep learning

- Recipe for success





# Discriminative deep learning

- Recipe for success:




Google's winning entry  
into the ImageNet 1K  
competition (with extra data).



# Discriminative deep learning


- Recipe for success:
  - Gradient backpropagation.
  - Dropout
  - Activation functions:
    - rectified linear
    - maxout

Google's winning entry  
into the ImageNet 1K  
competition (with extra data).



# Generative modeling

- Have training examples  $\mathbf{x} \sim p_{\text{data}}(\mathbf{x})$
- Want a model that can draw samples:  $\mathbf{x} \sim p_{\text{model}}(\mathbf{x})$
- Where  $p_{\text{model}} \approx p_{\text{data}}$



# Why generative models?

- Conditional generative models
  - Speech synthesis: Text  $\Rightarrow$  Speech
  - Machine Translation: French  $\Rightarrow$  English
    - French: Si mon tonton tond ton tonton, ton tonton sera tondu.
    - English: If my uncle shaves your uncle, your uncle will be shaved
  - Image  $\Rightarrow$  Image segmentation
- Environment simulator
  - Reinforcement learning
  - Planning
- Leverage unlabeled data

# Maximum likelihood: the dominant approach

- ML objective function

$$\theta^* = \max_{\theta} \frac{1}{m} \sum_{i=1}^m \log p(x^{(i)}; \theta)$$


# Undirected graphical models

- State-of-the-art general purpose undirected graphical model: **Deep Boltzmann machines**
- Several “hidden layers”  $h$

$$p(h, x) = \frac{1}{Z} \tilde{p}(h, x)$$

$$\tilde{p}(h, x) = \exp(-E(h, x))$$


$$Z = \sum_{h,x} \tilde{p}(h, x)$$



# Undirected graphical models: disadvantage

- ML Learning requires that we draw samples:

$$\frac{d}{d\theta_i} \log p(x) = \frac{d}{d\theta_i} \left[ \log \sum_h \tilde{p}(h, x) - \log Z(\theta) \right]$$




- Common way to do this is via MCMC (Gibbs sampling).

# Boltzmann Machines: disadvantage

- Model is badly parameterized for learning high quality samples.
- Why?
  - Learning leads to large values of the model parameters.
    - ▶ Large valued parameters = peaky distribution.
  - Large valued parameters means slow mixing of sampler.
  - Slow mixing means that the gradient updates are correlated  $\Rightarrow$  leads to divergence of learning.


# Boltzmann Machines: disadvantage

- Model is badly parameterized for learning high quality samples.
- Why poor mixing?



Coordinated  
flipping of low-  
level features

MNIST dataset




1st layer features (RBM)

# Directed graphical models

$$p(x, h) = p(x \mid h^{(1)})p(h^{(1)} \mid h^{(2)}) \dots p(h^{(L-1)} \mid h^{(L)})p(h^{(L)})$$

$$\frac{d}{d\theta_i} \log p(x) = \frac{1}{p(x)} \frac{d}{d\theta_i} p(x)$$

$$p(x) = \sum_h p(x \mid h)p(h)$$




- Two problems:
  1. Summation over exponentially many states in  $\mathbf{h}$
  2. Posterior inference, i.e. calculating  $p(\mathbf{h} \mid \mathbf{x})$ , is intractable.

# Directed graphical models: New approaches

- The Variational Autoencoder model:
  - Kingma and Welling, *Auto-Encoding Variational Bayes*, International Conference on Learning Representations (ICLR) 2014.
  - Rezende, Mohamed and Wierstra, *Stochastic back-propagation and variational inference in deep latent Gaussian models*. ArXiv.
  - Use a reparametrization that allows them to train very efficiently with gradient backpropagation.

# Generative stochastic networks

- **General strategy:** Do not write a formula for  $p(x)$ , just learn to sample incrementally.



- **Main issue:** Subject to some of the same constraints on mixing as undirected graphical models.

# Generative adversarial networks

- Don't write a formula for  $p(x)$ , just learn to sample directly.
- No summation over all states.
- How? By playing a game.

# Two-player zero-sum game

- Your winnings + your opponent's winnings = 0
- Minimax theorem: a rational strategy exists for all such finite games

# Two-player zero-sum game

- Strategy: specification of which moves you make in which circumstances.
- Equilibrium: each player's strategy is the best possible for their opponent's strategy.
- Example: Rock-paper-scissors:
  - Mixed strategy equilibrium
  - Choose your action at random

			Your opponent		
			Rock	Paper	Scissors
			Rock	0	-1
			Paper	1	0
			Scissors	-1	1


# Generative modeling with game theory?

- Can we design a game with a mixed-strategy equilibrium that forces one player to learn to generate from the data distribution?

# Adversarial nets framework

- A game between two players:
  1. Discriminator D
  2. Generator G
- D tries to discriminate between:
  - A sample from the data distribution.
  - And a sample from the generator G.
- G tries to “trick” D by generating samples that are hard for D to distinguish from data.

# Adversarial nets framework



# Zero-sum game

- Minimax objective function:

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

- In practice, to estimate  $G$  we use:

$$\max_G \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log D(G(\mathbf{z}))]$$


Why? Stronger gradient for  $G$  when  $D$  is very good.

# Discriminator strategy


- Optimal strategy for any  $p_{\text{model}}(x)$  is always

$$D(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_{\text{model}}(x)}$$


# Learning process




# Learning process



# Learning process



# Learning process



# Theoretical properties

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

- Theoretical properties (assuming infinite data, infinite model capacity, direct updating of generator's distribution):
  - Unique global optimum.
  - Optimum corresponds to data distribution.
  - Convergence to optimum guaranteed.

# Quantitative likelihood results

- Parzen window-based log-likelihood estimates.
  - Density estimate with Gaussian kernels centered on the samples drawn from the model.

Model	MNIST	TFD
DBN [3]	$138 \pm 2$	$1909 \pm 66$
Stacked CAE [3]	$121 \pm 1.6$	$2110 \pm 50$
Deep GSN [6]	$214 \pm 1.1$	$1890 \pm 29$
Adversarial nets	$225 \pm 2$	$2057 \pm 26$


# Visualization of model samples




MNIST



TFD




CIFAR-10 (fully connected)




CIFAR-10 (convolutional)

# Learned 2-D manifold of MNIST




# Visualizing trajectories

1. Draw sample (A)
2. Draw sample (B)
3. Simulate samples along the path between A and B
4. Repeat steps 1-3 as desired.



# Visualization of model trajectories




MNIST digit dataset



Toronto Face Dataset (TFD)

# Visualization of model trajectories

CIFAR-10  
(convolutional)



# Extensions

- Conditional model:
  - Learn  $p(x | y)$
  - Discriminator is trained on  $(x, y)$  pairs
  - Generator net gets  $y$  and  $z$  as input
  - Useful for: Translation, speech synth, image segmentation.

# Extensions

- Inference net:
  - Learn a network to model  $p(z | x)$
  - Infinite training set!

# Extensions

- Take advantage of high amounts of unlabeled data using the generator.
- Train  $G$  on a large, unlabeled dataset
- Train  $G'$  to learn  $p(z|x)$  on an infinite training set
- Add a layer on top of  $G'$ , train on a small labeled training set

# Extensions

- Take advantage of unlabeled data using the discriminator
- Train G and D on a large amount of unlabeled data
  - Replace the last layer of D
  - Continue training D on a small amount of labeled data

Thank You.

Questions?